

Predicting Cross-Sectional Relative Returns in the S&P 500: A Feature-Engineering and Walk-Forward Modeling Case Study

Quantitative Research Case Study

1 Problem and Approach

The task is to predict *relative* future stock returns inside the S&P 500: not whether the market goes up, but which stocks will out- or under-perform their peers. This framing is deliberate. Index-level direction is dominated by macro factors that daily OHLCV data cannot forecast, whereas the *cross-section* of returns carries persistent, tradable structure (momentum, reversal, liquidity, volatility, and sector-rotation effects) that is learnable from price and volume alone.

We therefore predict each stock’s forward return *in excess of its market and sector*, first test validation-selected model ensembles, and then use the same machinery to build a narrower momentum-plus-residual overlay. The primary trading evaluation is a cost-adjusted, sector-neutral decile long–short portfolio: within each sector, long the top-scored stocks and short the bottom-scored stocks, then combine sectors into a dollar-neutral book. The entire pipeline is built to be honest about look-ahead leakage, overlapping-window inflation, transaction costs, and survivorship bias, because in this problem the easiest way to produce a spectacular-looking result is to make a methodological mistake.

Headline design. Curated technical/liquidity/cross-sectional features → sector- and market-relative forward-return targets → a single-factor momentum baseline → compact Ridge/XGBoost models trained on factor-neutral residual labels → a validation-selected momentum-plus-residual overlay → rank IC / ICIR, DSR/PBO, style-neutral attribution, capacity, and cost-aware portfolio backtests. The research decision is simple: on price/volume plus sector metadata alone, trade or benchmark against *t*-turnover 12–1 momentum at the shorter traded horizons where the hold-out evidence is strongest, do not promote the ML/graph/residual models as standalone alpha, and spend the next research dollar on a true point-in-time universe with delisting returns. The most statistically reliable tradable benchmark is the 5-day sector-neutral momentum book: it has lower Sharpe than the 20-day point estimate, but it is the only traded row that combines a positive block-bootstrap IC interval with the lowest turnover (0.22). Its net Sharpe is 0.63, but the portfolio estimate itself is imprecise (mean-return $t = 1.31$, bootstrap Sharpe CI $[-0.32, 1.63]$), so the endorsement rests on cross-sectional IC and implementation cost rather than on a significant standalone portfolio Sharpe. The endorsed signal is materially survivorship-sensitive: the partial inclusion-side PIT filter cuts the 30-day momentum Sharpe from 0.741 to 0.446. The headline result is therefore negative and deflation-first: on this current-membership S&P 500 dataset, no broad multivariate, graph, MLP, or residual overlay construction produces statistically deflated alpha beyond the 12–1 momentum benchmark. The 30-trading-day sector-neutral residual overlay reduces volatility and drawdown versus standalone momentum, but it does not add statistically distinguishable return: over 35 non-overlapping test periods the paired return-difference *t*-stat is only 0.25 and the overlay is 93% correlated with momentum. Its Rank IC gain is negligible, the bootstrap Sharpe gap interval crosses zero, and the strict full-grid DSR/PBO audit rejects a production alpha interpretation. Momentum’s status as the core benchmark is anchored mainly on the pre-registered 1-day cross-sectional IC confirmation (plain $t = 3.54$; Newey–West lag-21 $t = 4.39$) and the external momentum literature, not on pretending that one 35-period 30-day portfolio Sharpe is precise. At traded horizons the same signal is lower-power: a focused HAC/block check gives hold-out block-bootstrap t of 2.41/2.13/1.89/1.85 at 5/10/20/30 days, with the 20d/30d block-bootstrap IC intervals crossing zero. The standalone 20-day momentum book has a higher Sharpe point estimate (0.82), but its IC interval crosses zero and its portfolio Sharpe is still imprecise (mean-return $t = 1.69$, bootstrap CI $[-0.12, 1.90]$). The standalone 30-day momentum book has Sharpe 0.740, mean-return $t = 1.51$, and an i.i.d. bootstrap Sharpe interval of $[-0.22, 1.85]$.

2 Data

We use only two files: `sp500_stocks.csv` (daily OHLCV per ticker) and `sp500_companies.csv` (sector, sub-industry, headquarters, founding year, and index `date_added`). No fundamentals, analyst, macro, or external listing data are used. The cleaned panel covers 503 symbols from 2000-01 to 2026-06 (2.926M stock-day rows).

Price adjustment (data quality). The raw file exposes only a `close` column and no separate `adj_close`. We therefore audited major split events and verified that this `close` behaves as an adjusted close price: around NVDA’s 2024-06-10 10:1 split (120.68 → 121.58) and AAPL’s 2020-08-31 4:1 split (121.06 → 125.17) there is no split-day gap, consistent with a `yfinance auto_adjust=True` export. In this study, `close` is therefore treated as the adjusted close used for all returns, momentum, volatility, and forward-return targets. Residual uncertainty about dividend adjustment is treated as a minor data-quality caveat.

Raw cleaning. The CSV-to-Parquet preparation step normalizes ticker strings, enforces one row per `symbol,date`, drops rows with missing core OHLCV fields, non-positive prices or volume, and impossible OHLC relationships (`high` below open/low/close or `low` above open/high/close). The raw stock file has no core missing values or duplicate stock-days, but it does contain 4,413 zero-volume rows and one OHLC inversion; these 4,414 rows are removed before feature construction. Large but plausible moves are audited rather than deleted: 44 stock-days have absolute one-day returns above 50%, while none exceed the 300% hard error threshold. Company metadata are also normalized and deduplicated; all 503 tickers have sector and sub-industry labels.

Survivorship bias. The constituent list is effectively the *current* membership applied to history, so firms that went bankrupt, were delisted, or were acquired (Enron, Lehman, Bear Stearns, WaMu, Sears, Kodak, Yahoo, Celgene, Xilinx, ...) are absent. This is a major current-membership bias, although the sign for a long–short momentum sleeve is not identified without delisting returns and side-specific holdings. The metadata include `date_added`, which supports a partial point-in-time inclusion check: 72 current

Table 1: Single-factor sanity check on core features, target `target_excess_sector_fwd_5d`. Mean IC is the average daily Spearman rank correlation after the production cleaning/evaluation controls.

Feature	Mean IC	Interpretation
<code>sector_rank_sma_gap_20d</code>	-0.0131	within-sector short-term reversal
<code>ret_5d</code>	-0.0125	recent winners partially mean-revert
<code>amihud_20d</code>	+0.0123	liquidity / price-impact information
<code>log_dollar_volume</code>	-0.0123	liquidity/size relation in relative returns
<code>mom_252d_skip_21d</code>	+0.0102	12–1 month momentum survives weakly
<code>rsi_14d</code>	-0.0099	overbought names tend to cool off

members were added on or after the 2022 hold-out start. We use that field as a sensitivity in §6, but still do not claim a true survivorship correction because deleted names and delisting returns are missing.

3 Feature Engineering

Each row is one stock-day observation, and every feature uses only information available at or before that date (all rolling windows are backward-looking with an explicit `shift`). We build ≈ 500 features in economically motivated groups:

- **Price / momentum / reversal:** multi-horizon returns (1–252d), log returns, skip-window momentum, momentum acceleration, short-term reversal, intraday/overnight decomposition, close-in-bar location. *Captures short-term reversal and medium-term momentum.*
- **Volatility / tail / drawdown:** rolling, upside/downside, Parkinson and Garman–Klass volatility, ATR, skew/kurtosis, drawdown from recent highs. *Controls for risk state and crash sensitivity.*
- **Trend / technical:** SMA/EMA gaps and slopes, MACD, RSI, Bollinger z-score/bandwidth/%b, channel position, breakout distance, Williams %R, CCI. *Encodes overbought/oversold and trend regimes.*
- **Liquidity / volume:** log dollar volume, volume z-scores and momentum, Amihud illiquidity, return-per-dollar-volume. *Captures attention, liquidity shocks, and transaction-cost pressure.*
- **Market regime:** equal-weight market return, breadth, cross-sectional dispersion, market volatility. *Relative-return predictability is regime-dependent.*
- **Statistical linkage / residual risk:** rolling market beta, market correlation, idiosyncratic return, and idiosyncratic volatility. *Captures stock-level residual behavior after market exposure is removed.*
- **Sector / sub-industry and peer / style / geography:** group returns, breadth, dispersion, leave-one-out peer means, headquarters-region peers, and style buckets. *Separates stock-specific alpha from group movement.*
- **Graph relation features:** two graph feature families are kept separate. The original `fixed_graph` variant uses deterministic same-sector/sub-industry, style-nearest-neighbor, and rolling-correlation graph relation embeddings. The `supervised_graph` variant uses out-of-fold supervised GAT embeddings and scores. *This separates static peer-structure features from supervised relation-aware graph learning.*
- **Cross-sectional transforms:** date-wise z-scores/ranks plus sector- and sub-industry-neutral versions of the key signals. *Because the target is relative, cross-sectional normalization is usually more informative than raw levels.*

Feature audit and sanity check. The rebuilt feature manifest contains 525 non-target columns and 40 target columns. The large count is not 525 independent ideas; it comes from expanding a smaller set of economic signals across lookback windows, related definitions, and neutralization transforms. For example, the cross-sectional layer alone contributes 168 columns by applying date-wise, sector-wise, and sub-industry-wise z-scores and ranks to selected base signals. The default walk-forward models do not use the full catalog: they use a curated 61-column `core` set covering momentum/reversal, volatility, technical indicators, liquidity, market/industry context, and key cross-sectional ranks. The executable grid treats graph inclusion as a feature variant: `tabular` uses those 61 core features, while `fixed_graph` (called `graph` in some lower-level pipeline outputs) appends 19 deterministic graph columns (`graph_emb_0-graph_emb_15` plus relation-fusion weights for sector, style kNN, and rolling correlation). The three relation weights (`graph_rel_weight_sector`, `graph_rel_weight_style_knn`, and `graph_rel_weight_rolling_corr`) are normalized across relations and therefore sum to about one up to floating-point error; for each stock-date they indicate whether the graph relation embedding relies more on sector peers, style-nearest peers, or rolling-correlation peers. Separately, `supervised_graph` appends target-matched out-of-fold GAT score, embedding, and relation-weight columns, while `graph_supervised` is the combined fixed-graph plus supervised-GAT variant. All variants use the same targets, folds, models, costs, purge/embargo, and hold-out protocol.

We explicitly verified that the expected factor families are present: 12-minus-1-month momentum (`mom_252d_skip_21d`), short-term reversal (`rev_*`), volatility/range estimators, Bollinger features, MACD, RSI-style oscillators, Amihud illiquidity, and date/sector/sub-industry z-score and rank transforms. The short-cycle feature block also includes the signals most likely to be orthogonal to 12–1 momentum: 1–5 day reversal, Amihud liquidity/impact, volume shocks, market breadth, cross-sectional dispersion, and rolling idiosyncratic return after market-beta removal. As a usefulness check, we ran a single-factor daily rank-IC scan on the cleaned 2008–2026 panel (after long rolling-window warm-up for 12–1 momentum and linkage features) using the same T+1 label shift, bottom-10% liquidity filter, and 1%/99% within-date winsorization as the modeling pipeline. Table 1 shows the largest effects in the core set. The magnitudes are small, as expected for daily cross-sectional equity alpha, but the signs are economically coherent: short-term overextension tends to reverse, 12–1 momentum is positive, and liquidity/impact variables contain weak but measurable information.

Prediction target. For horizon h we define the forward total return $r_{i,t}^{(h)}$ and subtract a contemporaneous group mean to obtain

sector/market-relative targets, e.g. $\text{target_excess_sector_fwd_hd}_{i,t} = r_{i,t}^{(h)} - \bar{r}_{\text{sector}(i),t}^{(h)}$. Targets are forward-looking and are strictly excluded from the feature matrix. The target generator also writes a date-wise percentile-rank target, `target_rank_fwd_hd`, so short-horizon experiments can use volatility-normalized, ranked, or winsorized labels when raw returns are too noisy. The implemented horizon grid covers $h \in \{1, 5, 10, 20, 30, 40, 50, 60, 70, 80, 90\}$. The original broad search emphasized 5d/10d, but the final Route B rerun uses the 30d sector-relative horizon because the shorter horizons did not deliver reliable net economics after costs. This is a selection step, not an exogenous design fact, and is one reason the final overlay must pay the full multiple-testing charge.

Execution timing. The raw target columns are generated on each feature date, but the model is evaluated with a one-day implementation lag. A feature row at date t is interpreted as information available after the t close; the signal is therefore tradeable only from $t+1$. In code, the effective training label is the raw target shifted one trading day forward, so a row dated t is matched to the forward return window that starts from the next trading day. This is the convention used for both rank-IC evaluation and portfolio PnL.

4 Methodology

4.1 Models

We compare, from simple to flexible: (i) a **single-factor momentum** baseline (rank by 12-minus-1-month momentum); (ii) **Ridge** regression (linear, strong regularization); (iii) **LightGBM** and (iv) **XGBoost** gradient-boosted trees (non-linear, interaction-aware); and (v) a small **PyTorch MLP** (two hidden layers, dropout, weight decay, early stopping on validation). The MLP is kept deliberately small and is *not* tuned until it wins; §5 reports its honest standing — competitive on in-sample validation IC, but not stable enough to headline once the untouched hold-out, turnover, and momentum benchmark are applied.

Momentum-residual overlay. Because the single-factor momentum baseline is strong, the implemented post-training evaluator does not force ML to replace momentum. For each date, both the baseline score and model score are converted to cross-sectional z-scores, the model score is residualized on the momentum score, and the final overlay signal is

$$s_{i,t} = z(\text{momentum}_{i,t}) + \lambda z(\widehat{\varepsilon}_{i,t}),$$

where $\widehat{\varepsilon}_{i,t}$ is the model component orthogonal to momentum within that day’s cross-section. The overlay weight is selected from a small pre-registered grid on validation net portfolio metrics, with $\lambda=0$ reproducing the momentum baseline. The final Route B run selects $\lambda=0.2$ on validation Sharpe; the test split is not used to tune the amount of residual ML alpha.

Route B residual-target training. After the broad model/graph search failed to produce a robust replacement for momentum, I reran the main experiment with a narrower design: each effective forward label is residualized cross-sectionally by date against sector dummies and OHLCV style proxies (liquidity, volatility, beta/idiosyncratic volatility, short-term reversal, 12–1 momentum, overnight/intraday return, volume shock, breadth and dispersion). Ridge/XGBoost then train on this residual target. The final tradable signal is still evaluated on raw forward returns and is combined with the momentum baseline as a small residual sleeve rather than a replacement.

Graph feature variants. Graph features are predictive inputs, not separate trading rules. The grid distinguishes deterministic `fixed_graph` peer embeddings from out-of-fold supervised-GAT embeddings and their combined variant, all under the same purge, embargo, execution-lag, validation, cost, and portfolio rules. The results below treat graph learning as an ablation family only: it does not earn a headline alpha claim.

4.2 Validation-selected rank ensemble

The ensemble is the final *prediction layer*, not a portfolio construction rule. For each base model m and date t , raw predictions are converted into a cross-sectional percentile rank $q_{i,t}^{(m)} = \text{rank}_t(\hat{y}_{i,t}^{(m)})$. The final score is then an average of member ranks,

$$\text{final_score}_{i,t} = \sum_{m \in \mathcal{M}} w_m q_{i,t}^{(m)},$$

with either equal weights or non-negative weights proportional to validation rank ICIR. Member selection uses validation rank ICIR and stability gates rather than validation mean IC; the post-validation test split is never used to select models, tune hyperparameters, choose ensemble members, or choose weights.

We compare five ensemble definitions:

- **ensemble_all_val_selected:** the current full ensemble, using the validation-selected best run for each model and feature variant.
- **ensemble_no_mlp:** the same rule after excluding all MLP members.
- **ensemble_tree_only:** Ridge, LightGBM, and XGBoost members only.
- **ensemble_supervised_graph_tree:** only supervised-graph Ridge, LightGBM, and XGBoost members.
- **ensemble_stability_filtered:** only members with positive validation rank IC, enough positive validation years, and non-negative available validation-regime rank IC.

This is also why the robustness tables report **ensemble_no_mlp** and **ensemble_tree_only**: the triage runs show that MLPs can look good on validation IC but are not stable enough to headline unless they pass the same walk-forward and regime gates.

4.3 Leakage controls

Every result below is produced under the same defenses:

- **One-day execution lag:** scores computed on day t are traded at $t+1$, so no same-bar information is used.
- **Label purge:** training rows whose forward label window crosses a split boundary are dropped, on both the training and the evaluation side.
- **Horizon embargo:** the validation/test window starts an h -day embargo after training ends, removing label overlap at the seam.
- **Train-only fitting:** imputation medians, scaling statistics, and the model are fit on training rows only.
- **Universe & winsorization:** before splitting, each date’s cross-section drops the bottom 10% by dollar volume and clips features at the 1%/99% within-date quantiles, so illiquid names and outliers do not drive the fit or an untradeable backtest.

4.4 Avoiding overlapping-window inflation

A subtle but decisive issue: if an h -day forward return is booked every day, consecutive observations overlap by $h-1$ days, their IC series is strongly autocorrelated, and a naive $\text{ICIR} = \sqrt{252} \cdot \text{IC} / \sigma(\text{IC})$ together with an overlap-booked Sharpe explode with h — a pure artifact, not alpha. We fix this in two ways: (1) the portfolio rebalances every h days so booked returns are *non-overlapping*; and (2) ICIR is annualized from a non-overlapping IC subsample (one observation per h days, scaled by $\sqrt{252/h}$). We retain the naive value as `rank_ic_ir_raw` only to document the inflation (Fig. 5). A sanity gate flags any run with $\text{Sharpe} > 3$ or $\text{ICIR} > 5$ as suspected leakage/overlap.

4.5 Validation: walk-forward selection + hold-out

A single static split can land in a lucky regime. The single-model horizon comparison therefore uses **expanding yearly walk-forward**: train on all data up to year Y , validate on year $Y+1$ (with purge + embargo), then roll forward one year and expand the training window. Model selection uses *walk-forward validation ICIR* and fold stability, not validation mean IC. The fixed 61-feature core is used instead of the full 525-column catalog for the same reason: it retains the economically interpretable, stable feature families while avoiding a high-variance short-horizon fit. The larger supervised-graph ensemble grid uses the same leakage controls but a fixed research split for tractability: train through 2018, validate through 2021, and evaluate the post-validation 2022–2026 test split once. In both cases, test results never participate in feature, model, hyperparameter, ensemble-member, or weight selection.

4.6 Portfolio and metrics

On each rebalance date we sort stocks by the reported score, either the validation-selected ensemble from §4.2 or the final Route B momentum-plus-residual overlay. Portfolio construction is deliberately written as a ladder from simple to stringent:

- **Long-only top bucket:** hold the top decile, either equal-weighted or score-weighted, with no short book. This is closest to a constrained long-only product and is evaluated with annualized return, volatility, Sharpe, max drawdown, turnover, and cost-adjusted return.
- **Plain long–short decile:** long the top 10%, short the bottom 10%, with equal dollar capital on both sides. This is the standard alpha ranking test corresponding to rank IC.
- **Sector-neutral long–short decile:** rank stocks within each sector, build top-minus-bottom decile books inside usable sectors, and average sector books into the total portfolio. This is the primary portfolio in the reported ensemble backtest because the target itself is sector-relative.
- **Market/sector beta-neutral portfolio:** an institutional-style robustness extension that constrains portfolio beta and sector exposures near zero in addition to dollar neutrality. It is stricter than the decile rule, but requires an optimizer and typically reduces raw return.
- **Cost-aware / turnover-controlled variants:** deduct explicit transaction costs, rebalance every 5/10/20 trading days or every target horizon, smooth scores with EWMA, and optionally restrict rebalancing to material signal changes through a no-trade band. The residual-turnover evaluator also sweeps no-trade bands of 0/5/10/15% of book turnover, which targets the gap between model turnover (~ 0.55 – 0.70 per rebalance) and the slower momentum benchmark (~ 0.16).

The empirical tables below use the sector-neutral long–short decile as the primary portfolio, rebalanced every h days with a 5 bps per-side transaction cost applied to realized turnover. We report rank IC, overlap-adjusted ICIR, decile spread, annualized net return, volatility, net Sharpe, max drawdown, and turnover.

4.7 Frequency-band and sleeve-book diagnostics

The final diagnostic pass asks whether a product should remain model-light: instead of treating LGBM/XGBoost as the center of the system, test whether simple economic sleeves survive their own native horizons and whether ML earns a small, explicitly deflated role. LGBM is therefore not needed for the core factor book; it is only a falsifiable add-on for regime-conditional sleeve weights, within-band residual combination, or a shrinkage portfolio optimizer. Each of those roles is charged with the same DSR/CPCV skepticism as the model grid, because the available regime episodes are few and the factor-level Fama–MacBeth evidence is weak. The signal-frequency diagnostic pre-registers 12 single factors and evaluates them against sector-relative labels at horizons 1, 5, 20, 40, 60, and 90 days. It reuses the production T+1 execution lag, bottom 10% liquidity filter, 1%/99% within-date winsorization, and overlap-adjusted ICIR. For each signal, the native horizon is selected by maximum absolute discovery ICIR using data through

2021; the 2022+ hold-out confirms or rejects that band. ICs are computed by re-ranking both signal and target on the common non-missing cross-section for each date.

The two-sleeve book then uses only the signals that survived the hold-out sign screen: 12–1 momentum and short reversal. For each sleeve k and rebalance date t , sector-neutral decile weights $w_{i,t}^{(k)}$ are formed from the signed raw signal. Combined weights are constructed at the position level, not by blending return streams,

$$w_{i,t}^{\text{book}} = \sum_k \alpha_k w_{i,t}^{(k)}, \quad \alpha_k = \frac{1/\hat{\sigma}_{k,\text{disc}}}{\sum_j 1/\hat{\sigma}_{j,\text{disc}}},$$

where $\hat{\sigma}_{k,\text{disc}}$ is the discovery-period net-return volatility. Because each sleeve is dollar-neutral, the combined book remains dollar-neutral. Costs are charged from the actual combined position turnover $\tau_t = \frac{1}{2} \sum_i |w_{i,t}^{\text{book}} - w_{i,t-1}^{\text{book}}|$.

For the multiple-testing check, the combined book’s hold-out Sharpe is deflated in per-observation Sharpe units. With per-rebalance Sharpe SR , skew γ_3 , kurtosis γ_4 , and T observations, the PSR denominator is

$$d = \sqrt{1 - \gamma_3 SR + \frac{\gamma_4 - 1}{4} SR^2}.$$

The expected maximum null Sharpe for N searched single-factor trials is

$$SR^* = \sigma_{SR} \left[(1 - c_E) \Phi^{-1} \left(1 - \frac{1}{N} \right) + c_E \Phi^{-1} \left(1 - \frac{1}{N_e} \right) \right],$$

where $c_E = 0.5772$ is Euler’s constant and σ_{SR} is the cross-trial standard deviation of per-observation Sharpes. The reported DSR is

$$\Phi \left(\frac{(SR - SR^*) \sqrt{T-1}}{d} \right),$$

with $N = 12$ for the pre-registered single-factor scan. The Route B full-grid DSR uses the same formula with $N = 884$, counting the archived searched model/target/horizon/configuration rows in the mirrored walk-forward and residual-overlay experiment grid; this is intentionally conservative because the rows pool heterogeneous horizons and target families.

5 Results

Validation selection. Selecting by across-fold mean validation ICIR, the small MLP is the strongest model at the 1-, 5-, and 30-day horizons (e.g. sector-relative 5d ICIR 1.15, market-relative 0.99) and LightGBM at 20 days (market-relative 0.99); Ridge and XGBoost trail (Fig. 4). No single model dominates every horizon and the boosted trees do *not* uniformly win — the MLP is genuinely competitive on the in-sample IC metric, so we do not pretend trees are categorically better.

Ensemble and graph comparison. The 5-day rank ensembles do not change the investment conclusion. Excluding MLP improves the post-validation rank IC relative to the full ensemble, and the supervised-graph tree ensemble is the strongest 5-day ranker among the archived ensemble diagnostics, but net Sharpe is only about zero to 0.05 after costs and turnover remains near 0.7 per rebalance. I therefore keep ensemble/GAT results as negative model-selection evidence rather than as a tradable alpha claim.

The decisive test: the untouched hold-out. The validation ranking does not survive out of sample. On the 2022–2026 hold-out the *single-factor momentum baseline* (rank by 12-minus-1-month momentum, no fitting) beats every trained multivariate model at the original 5-day benchmark horizon on rank IC (0.023 vs ≤ 0.013), ICIR (1.02 vs ≤ 0.73), and turnover (0.22 vs ~ 0.59). Its corrected 5-day Sharpe is 0.63, not the legacy > 1 figure, so the honest claim is narrower: momentum is the cleaner rank signal and the lower-turnover book, while Sharpe leadership depends on horizon and evaluator convention (Table 5). The models’ validation-IC edge is therefore largely overfit: a robust, low-turnover factor is hard to beat after costs. This is the central, deliberately humbling finding of the study.

Signal-frequency diagnostic. The single-factor frequency scan provides an independent check on that conclusion. The discovery-period ICIR surface shows economically sensible bands: short reversal signals are negative at 1–5 days, low-risk variables flip sign from short to longer horizons, and Amihud illiquidity is the discovery-period champion across 20–60 days, but the 2022+ confirmation step is much harsher (Table 2). At its native 1-day confirmation horizon, momentum is the only signal with a positive, significant, and stronger hold-out t -stat. Amihud’s in-sample strength collapses from $t = 5.51$ to $t = 0.41$, volume shock flips sign, and reversal keeps the expected sign but is weak. A follow-up traded-horizon momentum check using `report/artifacts/momentum_ic_traded_horizons/` keeps the horizon distinction explicit: block-bootstrap t is 2.41 at 5d, 2.13 at 10d, 1.89 at 20d, and 1.85 at 30d; the corresponding block-bootstrap mean-IC intervals are positive at 5d/10d but cross zero at 20d/30d. This is a clean warning that 1-day IC confirms signal existence, while tradability is horizon- and cost-limited even before adding ML. The daily Newey–West t -stat rising from the plain $t = 3.54$ to lag-21 $t = 4.39$ is not a HAC sign error; the daily IC series is mildly negatively autocorrelated, so the Bartlett long-run variance estimate tightens rather than widens.

Table 2: Pre-registered signal-frequency diagnostic. Native horizon is chosen by discovery-period absolute ICIR through 2021, then confirmed on the 2022+ hold-out under the production T+1, liquidity, winsorization, and overlap-adjusted ICIR rules.

Signal	Native h	Disc. ICIR	Disc. t	Test t	Hold-out read-through
mom_252d_skip_21d	1d	0.595	2.23	3.54	Only significant survivor; stronger OOS.
amihud_20d	40d	1.467	5.51	0.41	In-sample champion, OOS decay.
ret_5d reversal	1d	-1.027	-3.84	-0.94	Sign stable, but too weak.
Low-risk block	1d	-0.432 to -0.793	-1.62 to -2.97	≈ 0	Vol/idiosyncratic-vol/beta effect decays.
volume_z_20d	60d	0.514	1.93	-1.93	Sign flips on the hold-out.





Figure 2: Net equity curves for the 5-day momentum/reversal sleeve-book diagnostic.

Figure 1: Discovery signal $\times$ horizon ICIR surface for the pre-registered single-factor scan.

Route B 30-day residual overlay. The final research lead comes from the narrower residual-target rerun, not from the graph/MLP zoo. XGBoost is trained on 30-day sector-relative labels after daily ridge residualization against sector and OHLCV style proxies. The tradable score is then $z(\text{momentum}) + 0.2z(\text{residual XGBoost})$, with $\lambda = 0.2$ selected on validation Sharpe. Table 3 compares this Route B overlay with standalone momentum and the strongest plain 30d sector model on the 2022–2026 hold-out. This residual-label XGBoost rerun is separate from the broad-grid 30d sector MLP row in Table 5. It is a deliberately separate 30-day sleeve diagnostic, not the headline proof that momentum is tradable. The overlay improves the point-estimate economic profile relative to momentum (Sharpe 0.74 $\rightarrow$ 0.99, max drawdown $-13.8\% \rightarrow -7.3\%$), but a paired increment test shows that this is mostly volatility and drawdown dampening rather than new return: mean per-period return improves by only 0.06% (about 0.5% annualized), the paired return-difference t -stat is 0.25, the overlay–momentum correlation is 0.93, and a paired bootstrap Sharpe-gap CI is $[-0.17, 0.63]$. The evidence is therefore deliberately read conservatively: the sample has only $N = 35$ non-overlapping periods, the absolute bootstrap Sharpe CI is $[0.019, 1.796]$, full-grid DSR is 0.003 after 884 archived searched model/target/horizon/configuration trials, PBO is 50%, and the Rank IC change is within evaluator noise. The standalone momentum baseline is also thin-sample as a 30-day portfolio: Sharpe 0.740 has mean-return $t = 1.51$ and an i.i.d. bootstrap Sharpe interval of $[-0.22, 1.85]$ over the same 35 periods. The plain 30d LightGBM row shows why horizon framing matters: a raw model can beat the overlay’s Sharpe point estimate, but it fails the paper’s validation-selection rule (validation ICIR 0.0804 versus the selected MLP’s 0.4882), sits below the 884-row expected-max null Sharpe of 1.512, and lacks archived DSR/PBO artifacts. The overlay is kept only for its economic sleeve role and lower turnover, not for raw performance dominance.

Table 3: 30-day sector-relative hold-out comparison, net of 5 bps/side cost. Route B has $N = 35$ non-overlapping periods; the standalone momentum Sharpe bootstrap CI is $[-0.22, 1.85]$ and the overlay bootstrap Sharpe CI is $[0.019, 1.796]$. Full-grid DSR is 0.003 over 884 archived searched model/target/horizon/configuration trials (same-target DSR 0.814 over 14 trials), and CSCV PBO is 50% over 20 splits, so the overfit estimate is adverse but not precise. The model is a candidate residual sleeve, not a replacement for momentum; overlay IC/ICIR is reported under the same canonical raw-label convention as the standalone momentum row. The paired increment test uses archived returns in `report/artifacts/route_b_paired_increment/`; the standalone momentum bootstrap uses `report/artifacts/route_b_momentum_baseline/`. Annualized returns use arithmetic mean return times rebalance frequency; the paired artifact also stores geometric compounded figures.

Config	Source	Rank IC	ICIR	Net Sharpe	Ann. Ret	Max DD	Turnover
Momentum baseline	canonical baseline	0.0282	0.698	0.740	7.78%	-13.83%	0.401
Plain LightGBM model	not val-selected	0.0362	0.802	1.226	12.95%	-8.16%	0.631
Route B overlay, $\lambda = 0.2$	val-selected sleeve	0.0279	0.759	0.987	8.27%	-7.32%	0.494
Paired increment	overlay – momentum	–	–	$\Delta 0.246$	$\Delta 0.50\%$	$+6.51\text{pp}$	$\Delta 0.093$
Route B overlay, $\lambda = 0.3$	diagnostic only	0.0280	0.806	1.201	9.98%	-6.10%	0.521

Route B artifact check. The Route B model-search and robustness-audit outputs are mirrored locally under the Route B model-search and robustness audit output directories. A focused rerun of `eval_momentum_baseline.py` for the 30d sector-relative target reproduces the canonical 30d momentum economic row exactly: Sharpe 0.740, annual return 7.78%, max drawdown -13.83% , and turnover 0.401. I also recompute the selected Route B overlay score against the same raw forward label in `report/artifacts/route_b_canonical_ic/`: canonical Rank IC is 0.0279 versus momentum’s 0.0282, and ICIR is 0.759 versus 0.698. Under the residualized-label audit convention, the signs flip (overlay -0.0083 versus baseline -0.0105), which is useful only as a residual-target diagnostic. Both conventions say the same PM-relevant thing: the overlay’s rank-information gain is negligible; the economic comparison remains reproducible.

Route B paired increment. The archived paired return stream is the cleaner comparison than two absolute Sharpe estimates. The overlay and momentum returns are aligned across the same 35 rebalance dates; the overlay’s period volatility falls from 3.63% to 2.89%, but mean return improves by only 0.059% per period. The paired return-difference t -stat is 0.25, and the paired bootstrap

puts the Sharpe gap at +0.25 with a 95% interval $[-0.17, 0.63]$ and $P(\Delta \text{Sharpe} > 0) = 0.88$. This reframes the Route B lead as a high-correlation risk-sleeve candidate, not as independent residual alpha.

Momentum/reversal sleeve-book diagnostic. The next product-style check built the two-sleeve book implied by the frequency scan: momentum plus short reversal, rebalanced every 5 days and combined at the position level with discovery inverse-volatility weights. The estimated discovery weights are 43.1% momentum and 56.9% reversal, and the full-sample sleeve net-return correlation is mildly diversifying (-0.153), but the hold-out economics do not justify a second sleeve: momentum-only has the best 2022+ Sharpe (0.598), reversal-only is negative (-0.519), and the combined inverse-vol book has Sharpe 0.094 with DSR 0.055 after the 12-factor scan and bootstrap Sharpe interval $[-0.724, 1.095]$. This uses raw signed factor sleeves rather than the rank-transformed baseline in Table 5, so the Sharpe levels are not apples-to-apples; the useful conclusion is that a simple multi-sleeve product does not beat momentum in this dataset.

The full sleeve table is omitted for space; the relevant audit statistics are momentum-only Sharpe 0.598, reversal-only Sharpe -0.519 , inverse-vol combined Sharpe 0.094, DSR 0.055, and bootstrap Sharpe CI $[-0.724, 1.095]$. An earlier 10-day residual overlay diagnostic is likewise demoted to prose: it improved point-estimate Sharpe from 0.746 to 0.995 on the same hold-out, but the later Route B audit showed why such overlays must be reported as exploratory unless they survive deflation and attribution checks.

Graph relation ablation. The completed 120-candidate tree grid compares real relation families (sector, style-nearest-neighbor, rolling correlation, and all relations) against random-edge and no-edge controls for 5d/10d sector- and market-relative targets. The result is negative enough to summarize rather than headline: placebo controls beat real relations in three of four target blocks, and the one narrow real-relation win remains below the momentum ICIR. Graph learning is therefore treated as an experimental feature family, not a source of reported alpha.

Failed ablations and negative controls. The negative controls remain part of the research record but are compressed here because they are not alpha evidence: a 24-member fixed-graph tree seed ensemble reached only Rank IC 0.0171, ICIR 0.517, Sharpe 0.290, and turnover 0.649, while the best Kaggle/Ubiquant-style date-aggregate sector 5d test ICIR was 0.569, below the momentum ICIR of 1.016. These failures are important because they prevent the paper from presenting graph features, seed bagging, or date aggregates as successful complexity.

PM-level robustness audit. The Route B 30-day overlay is the best current residual-sleeve research lead, but it is still *not* a production-grade alpha claim. A post-training audit treats the full search as multiple testing, applies CSCV/PBO, neutralizes the final score against available style proxies, and stresses capacity and survivorship. The result is mixed: hold-out economics are better than momentum, and the same-target DSR is encouraging, but the full-grid DSR and PBO still force an exploratory interpretation (Table 4). The capacity, survivorship-haircut, and factor-neutral rows are sensitivity transforms of the same selected $N = 35$ overlay return stream; they bound fragility but do not provide independent confirmation.

Table 4: P0 robustness audit for the Route B 30-day residual momentum overlay. DSR is the probabilistic Sharpe ratio after deflating for the expected best result across the searched configurations; PBO is CSCV probability of backtest overfitting. DSR uses a pooled cross-trial Sharpe dispersion across heterogeneous searched horizons/targets, a conservative approximation rather than a claim that all trials are exchangeable.

Check	Result	Interpretation
Test Sharpe before deflation	0.987	Better than 30d momentum's 0.740, before selection correction.
Paired overlay — momentum test	$t = 0.25$; corr 0.93	Return increment is not statistically distinguishable; improvement is risk dampening.
Bootstrap Sharpe 95% CI	$[0.019, 1.796]$	Positive but wide.
DSR, same-target trials	0.814	Encouraging within the 14-trial 30d target search.
DSR, full grid trials	0.003	0.987 Sharpe is below the 884-row expected-max null of 1.512.
CSCV PBO	50.0%	High but low-precision overfit risk across 28 strategies and 20 splits.
Fama-MacBeth signal t -stat	0.05	No convincing independent residual-alpha coefficient.
Factor-neutral Sharpe, return, FM t	1.006; 8.27%→5.17%; 0.05	Turnover rises 0.49→0.70; no independent residual-alpha coefficient.
Capacity Sharpe at \$100M / \$1B	0.948 / 0.864	Capacity is plausible at modest scale; participation rises at \$1B.
PIT inclusion sensitivity	Mom. 0.741→0.446; overlay 0.987→0.725	Partial inclusion-side check weakens both rows; overlay remains exploratory.
Survivorship stress-haircut Sharpe	0.837	Sensitivity does not destroy the result, but remains uncorrected.




Figure 3: Best-model rank ICIR by horizon (validation vs hold-out), per target family.




Figure 4: Validation mean ICIR by model at the 5-day horizon.

Overlap correction matters. Fig. 5 contrasts the overlap-adjusted hold-out ICIR with the naive overlapping value. The naive metric inflates sharply with horizon (to ≈ 3 – 10 at 30–90 days) while the adjusted metric stays modest (≤ 1.3); reporting the naive number would manufacture a long-horizon “signal” that is pure window overlap. The $\text{Sharpe} > 3 / \text{ICIR} > 5$ sanity gate flags none of the adjusted runs.

Stability across regimes. Fig. 6 shows per-fold ICIR across the expanding walk-forward. The signal is positive but *variable* across folds, and the hold-out value (dashed) sits below the validation mean — the expected out-of-sample decay. The 90-day horizon is excluded from selection (too few non-overlapping validation periods for a stable ICIR), and the longer-horizon hold-out Sharpe values are estimated from few non-overlapping periods and should be read as noisy. The ensemble search also emits regime slices whenever a split overlaps the default stress windows (GFC 2008, COVID-2020, and inflation-2022); with the fixed ensemble split used here, the available ensemble stress rows are COVID-2020 in validation and inflation-2022 in test, while 2008 is covered only by the broader walk-forward stability analysis.




Figure 5: Overlap-adjusted vs naive (overlapping) hold-out ICIR by horizon.




Figure 6: Walk-forward per-fold ICIR for the short-horizon model-comparison run.

Hold-out performance. Table 5 reports the untouched post-2022 hold-out, net of 5 bps/side cost. The 5d and 20d blocks show the original conclusion: realistic cross-sectional equity “alpha” from price/volume features alone is small, and the trained models’ apparent in-sample edge does not translate cleanly into out-of-sample value. The added 30d block is included to avoid horizon-selective framing: plain 30d models have attractive point estimates, but their longer-horizon sample is small and they have not passed the Route B deflation/attribution audit.

6 Discussion

Overfitting and the validation–hold-out gap. The clearest evidence of overfitting is in the results themselves: models selected on validation ICIR fail to beat a single momentum factor on the original 5-day rank and turnover benchmark, and the model with the best validation ICIR (the MLP) is not the best economic hold-out model. The final 30-day Route B rerun improves the momentum book only after the ML component is constrained to a residual sleeve, and even then the paired increment test shows no statistically distinguishable return addition. Standalone residual models still have weak economics and high turnover. We use

Table 5: Untouched hold-out (2022–2026), net of 5 bps/side cost. The single-factor momentum baseline is sourced from `eval_momentum_baseline.py`; model rows are the walk-forward validation selections in `docs/horizon_comparison.md`. Parenthetical ICIR is the naive overlapping value, shown only to document inflation. The 5d/20d/30d blocks are all sector-relative to avoid target-family cherry-picking; the 20d momentum IC/ICIR comes from the traded-horizon robustness artifact, while the economic row is the same canonical momentum baseline evaluator. Bold is used only for the pre-30d benchmark blocks; the 30d block is too thin ($N = 35$ non-overlapping periods for Route B) and unaudited plain-model rows are not statistically rankable. Momentum portfolio uncertainty is in `report/artifacts/momentum_portfolio_uncertainty/`: 5d/20d/30d Sharpe CIs are $[-0.32, 1.63]$, $[-0.12, 1.90]$, and $[-0.22, 1.85]$.

Horizon	Config	Rank IC	ICIR (raw)	Net Sharpe	Ann. Ret	Turnover
5d	Momentum baseline (sector)	0.023	1.02 (2.05)	0.63	5.9%	0.22
5d	MLP (<i>val-selected, sector</i>)	0.013	0.73 (1.72)	0.17	1.2%	0.59
5d	Ridge (sector)	0.010	0.40 (0.99)	0.53	5.1%	0.65
5d	LightGBM (sector)	0.009	0.20 (1.17)	0.24	1.7%	0.53
20d	Momentum baseline (sector)	0.031	0.58 (–)	0.82	7.8%	0.36
20d	LightGBM (<i>val-selected, sector</i>)	0.021	0.48 (2.42)	0.65	5.3%	0.66
20d	Ridge (sector)	0.017	0.34 (1.66)	0.83	9.0%	0.60
30d	Momentum baseline (sector)	0.028	0.70 (3.04)	0.74	7.8%	0.40
30d	MLP (<i>broad-grid val-selected, sector</i>)	0.027	0.92 (3.88)	0.69	6.5%	0.70
30d	LightGBM (sector; not val-selected, < 1.512 null)	0.036	0.80 (4.34)	1.23	13.0%	0.63
30d	Route B overlay, $\lambda = 0.2$	0.028	0.76 (3.07)	0.99	8.3%	0.49

heavy regularization (Ridge penalty, tree depth/leaf/subsample limits, MLP dropout + weight decay + early stopping), select hyperparameters on validation folds only, and report a single untouched hold-out precisely so this gap is *visible* rather than hidden, with the Sharpe > 3/ICIR > 5 gate catching anything “too good”.

Multiple testing. The original grid is large enough that the best Sharpe must be treated as a selected maximum, not as a single pre-specified experiment. For the Route B 30-day overlay, raw hold-out Sharpe is 0.987 and the bootstrap 95% Sharpe interval is $[0.019, 1.796]$. The same-target DSR is 0.814, but full-grid DSR falls to 0.003 because the 884 archived searched model/target/horizon/configuration rows imply an expected-max null Sharpe of 1.512, above both the observed overlay Sharpe 0.987 and the visually stronger but non-selected LightGBM 30-day Sharpe 1.226. This DSR pools heterogeneous horizons and target families into one cross-trial null dispersion, which is a conservative penalty for the overlay, not a clean family-wise theorem. CSCV PBO is 50%, but only across 28 compact strategies and 20 splits, so it is an adverse instability flag rather than a precise probability. This is the most important statistical conclusion: the overlay is a materially better research lead than the old graph/MLP direction, but still not a statistically deflated production alpha. The paired overlay-minus-momentum test reinforces that conclusion: the overlay is 93% correlated with momentum, reduces period volatility, and has a return-difference t -stat of only 0.25. Symmetric rigor also weakens any overconfident reading of the benchmark portfolio itself: the 5-day, 20-day, and 30-day momentum Sharpe point estimates are positive but imprecise, with mean-return t -stats of 1.31, 1.69, and 1.51 and bootstrap Sharpe CIs that all cross zero. The separate factor-sleeve book is even weaker: after charging the 12 pre-registered single-factor trials, the combined momentum/reversal book has DSR 0.055 and a bootstrap Sharpe interval that crosses zero. It is useful as a falsification exercise, not as a new product claim.

Route B reproducibility. The synchronized Route B artifacts close the largest traceability gap: the 30d overlay metrics can now be traced from local CSV/parquet files rather than only from a remote run note. A focused local rerun of the 30d momentum baseline now matches the Route B baseline row, so the absolute economic comparison in Table 3 is reproducible under one canonical evaluator. The frequency-scan and sleeve-book diagnostics also write lightweight CSV snapshots under `report/artifacts/`. The reason to keep Route B exploratory is not provenance; it is weak independent attribution, wide confidence intervals, and the full-grid DSR/PBO charge.

Factor attribution. The dataset lacks a full commercial risk model and fundamental value/quality factors, but we can still neutralize the final signal against available style proxies: liquidity/size (`log_dollar_volume`), volatility (`vol_20d`), beta (`beta_60d`), Amihud illiquidity, short-term reversal (`ret_5d`), and 12–1 momentum. The Route B overlay’s score-neutralized Sharpe remains near 1.0, but that is not evidence of orthogonal alpha by itself: annual return falls from 8.27% to 5.17%, turnover rises from 0.49 to 0.70, and the non-overlapping Fama–MacBeth signal coefficient has only $t = 0.05$. The evidence therefore supports an economically interesting residual sleeve, not a clean claim of independent style-neutral alpha.

Look-ahead leakage. Addressed structurally: one-day execution lag, label purge on both sides, horizon embargo, and train-only preprocessing. We also verified that a no-embargo diagnostic inflates metrics, confirming the controls bind.

Survivorship bias. The dataset is current-membership-on-history, so the backtest is not point-in-time and cannot be treated as production evidence. The available `date_added` field supports an inclusion-side sensitivity: 72 current members were added on or after 2022-01-01, and after the baseline controls 61 such symbols appear in hold-out rows that a PIT inclusion filter removes. Recomputing the 30d sector-neutral momentum baseline with `sp500_member_asof==1` reduces Sharpe from 0.741 to 0.446, annual net return from 7.79% to 3.91%, and worsens max drawdown from -13.8% to -17.5%. Applying the same filter post-training to the selected Route B prediction file and rebuilding the overlay score reduces Sharpe from 0.987 to 0.725, annual net return from 8.27% to 5.69%, and worsens max drawdown from -7.3% to -10.0% (`report/artifacts/pit_inclusion_sensitivity/`). This is an important hit to the point estimates, not a new validation pass. The audit also adds a delisting-side parametric sensitivity: applying annual delisting rate $\times$ average delisting loss to half the book leaves the Route B overlay Sharpe at 0.951 in a base case (2% annual delisting rate, -30% average delisting return) and 0.837 in a stress case (5%, -50%). Neither exercise is a true correction. The 0.741→0.446 momentum haircut is therefore the survivorship-aware number a PM should keep in mind for this dataset, not a minor

footnote; the missing short side and delisting returns remain unidentified because absent names could have appeared on either side of the book. The right fix is point-in-time constituents and delisting returns.

Kronos and survivorship bias. Kronos is not used as a reported alpha model in this paper. Its relevant role is only as a possible future data-universe route: a sequence model trained on a point-in-time universe with delisted, acquired, and index-deleted names could generate price-volume embeddings beyond the current S&P 500 membership. I therefore treat Kronos as a candidate way to broaden the investable universe and reduce current-constituent survivorship bias, not as empirical evidence in the results above.

Data quality. Prices are split-adjusted; dividend adjustment is uncertain and treated as a minor caveat. We winsorize features and filter the illiquid decile so bad ticks and microcaps do not dominate.

Turnover, capacity, and market impact. We charge 5 bps per side on realized turnover and rebalance only every h days. Turnover is itself decisive: the trained residual models churn $\sim 75\text{--}85\%$ per rebalance, while the Route B overlay turns over 49.2% versus 40.1% for the 30-day momentum benchmark. A simple square-root impact stress test, calibrated to 10 bps at 1% ADV participation, gives overlay Sharpe 0.948 at \$100M AUM, 0.900 at \$500M, and 0.864 at \$1B, with 95th percentile participation rising to 25.3% at \$1B. This is not an execution model, but it is enough to show that capacity and short-side liquidity are first-order constraints.

Regime dependence. Walk-forward folds span 2008, 2020, and 2022 stress periods; Fig. 6 shows the signal is not uniform across regimes, motivating the market/breadth/dispersion regime features rather than a single static model. The frequency-band diagnostic makes the same point at the factor level: Amihud illiquidity is a strong discovery-period signal but dies in the 2022+ hold-out, while momentum is weak in parts of discovery and strong in the recent hold-out. This is why single-period IC surfaces are treated as hypothesis generators, not proof.

Economic intuition. The one signal that clearly survives out of sample is cross-sectional momentum — a well-documented behavioural underreaction effect — which is precisely why the single-factor baseline is so hard to beat. The sleeve-book diagnostic confirms that a theoretically attractive momentum-plus-reversal product is not enough by itself: reversal is mildly orthogonal, but too weak and too high-turnover to improve hold-out Sharpe. The Route B result is useful because it gives the model a narrower economic job: find residual structure that can be added to, not substituted for, momentum, and then survive DSR rather than merely improve a point estimate. The current Route B improvement is therefore best read as a candidate risk-timing effect, not as established residual alpha.

7 Limitations and Conclusion

This study delivers an honest, reproducible cross-sectional relative-return pipeline for the S&P 500 using only price/volume and light metadata. The key discipline is methodological: non-overlapping evaluation, walk-forward selection, explicit leakage controls, and a single-factor baseline that every model must clear in a clearly specified metric. The realistic conclusion is sobering: broad multivariate, graph, MLP, and residual-overlay searches do *not* pass a production-grade alpha audit. Momentum is the only clearly confirmed raw factor by 1-day IC confirmation and prior evidence, and the 5-day sector-neutral implementation is the most statistically reliable tradable benchmark in this study because its traded-horizon IC interval is positive and turnover is lowest (Sharpe 0.63, turnover 0.22, block-bootstrap IC $t = 2.41$). But this is not a production-ready endorsement: its own portfolio Sharpe is not statistically tight (mean-return $t = 1.31$, bootstrap Sharpe CI $[-0.32, 1.63]$), the partial PIT inclusion filter cuts the 30-day momentum Sharpe from 0.741 to 0.446, and the 30-day implementation is imprecise over only 35 non-overlapping periods ($t = 1.51$, bootstrap CI $[-0.22, 1.85]$). Its daily IC confirmation is more robust: a 21-day block bootstrap puts the hold-out mean Rank IC CI at $[0.0115, 0.0297]$ and Newey–West lag-21 gives $t = 4.39$; the traded-horizon HAC/block evidence is positive but only marginal by 20–30 days. Even the best Route B residual sleeve remains exploratory. Its 2022–2026 point estimate improves the 30-day momentum book’s Sharpe from 0.740 to 0.987 and drawdown from -13.8% to -7.3% , but the paired return increment is not statistically distinguishable ($t = 0.25$) and the overlay is 93% correlated with momentum. With canonical Rank IC 0.0279 versus momentum’s 0.0282, full-grid DSR 0.003, PBO 50%, and Fama–MacBeth $t = 0.05$, the result is best interpreted as volatility/drawdown dampening rather than established residual alpha. A later frequency-band and two-sleeve product diagnostic strengthens the conservative side of the conclusion: a momentum-plus-reversal inverse-vol book fails DSR. I would not trade the model as a standalone alpha, and I would not promote it without point-in-time constituents, delisting returns, one canonical portfolio evaluator, stronger style-factor attribution, and lower-turnover execution tests. The useful output is the triage: keep 12–1 momentum as the core benchmark, treat residual ML as a small research sleeve, and make DSR/PBO, PIT survivorship, capacity, and turnover part of model selection rather than post-hoc decoration.